

UNIVERSITY OF CALIFORNIA, DAVIS

ECS 158

PROGRAMMING ON PARALLEL ARCHITECTURES

---

# Parallelization of R's `combn()` Function

---

*Authors:*

Trisha FUNTANILLA  
Syeda INAMDAR  
Eva LI  
Jennifer WONG

*Professor:*

Norm MATLOFF

March 20, 2015

# Contents

# 1 Objective

Choose some function `g()` in R, either built-in to base R or in a CRAN package, to parallelize. Do so in Snow, OpenMP and CUDA (or substitute Thrust for either CUDA or OpenMP but not both), and run timing tests.

## 2 Function

We chose to parallelize the function `combn()` from the “combinat” package in CRAN.

### 2.1 Description

Taken from the official documentation:[?]

“Generate all combinations of the elements of `x` taken `m` at a time. If `x` is a positive integer, returns all combinations of the elements of `seq(x)` taken `m` at a time. If argument “`fun`” is not null, applies a function given by the argument to each point. If `simplify` is `FALSE`, returns a list; else returns a vector or an array. “...” are passed unchanged to function given by argument `fun`, if any.”

### 2.2 Why Speedups Are Possible

We have tried `combn()`, which has the inputs `combn(x, m, fun=NULL, simplify=TRUE, ...)`, with different inputs for the vector `x` and value `m`. We noticed that `combn()` became quite slow once `x` had a size numbering around 100 or higher and `m` around 5 or higher, taking several minutes to compute (but not print) its results. We think that we can achieve speedups by dividing the input `x` equally among the threads/nodes, and parallelizing the work of creating the output matrix or list. The goal is to speed up the function, which, based on its source code, currently computes combinations serially.

## 3 Parallelization

### 3.1 OpenMP

The C++ OpenMP program is called through R via `.Call()` using the Rcpp interface. It accepts the same input arguments as the original function, can handle the same types of valid inputs, and through new optional arguments, also allows the user to decide on a scheduling policy and chunk size.

Code highlights:

- Load balancing algorithm [?] (next section)
- Avoidance of critical sections and barriers

Other notes:

- Similar to the original function, the OpenMP implementation has a limit on the input/output size. The program terminates if R decides that it cannot allocate enough space for the output.

### 3.1.1 Load Balancing

Good load balancing is very critical to our program because for each element  $i$  in the input vector  $x$ , the number of combinations that can be generated with  $x_i$  is larger than that with  $x_{i+1}$ . A naive task distribution without proper load balancing could heavily skew the distribution of the workload for each thread. For example, directly assigning an  $x_i$  for each thread will result in the threads with lower id's generating far more combinations (thus, more work) than threads with higher id's).

In order to compensate for the load imbalance, the assignment of tasks to the threads was done in a wrap around manner. Using this method, for the first distribution, each thread will work on the element indexed in the same value as their id (e.g. thread 0 on  $x[0]$ , thread 1 on  $x[1]$ , and so on). Then, the following distribution depends on the total number of threads, the id of each thread, and the index of its previous assignment. With  $nth$  corresponding to the total number of threads and  $me$  corresponding to the thread id, the assignment following the initial distribution is determined using the equation:

$$next\_task = prev\_task + 2 * (nth - me - 1) + 1 \quad (1)$$

The distribution that will then follow the above depends on the id of each thread the index of the previous task. This next distribution is determined by the equation:

$$next\_task = prev\_task + 2 * me + 1 \quad (2)$$

The succeeding distributions alternates in using the two equations until the last distribution, which is for  $n - m + 1$  or the last  $x_i$  that can form a combination of size  $m$  with the elements succeeding it.

In this algorithm, each thread knows which indexes in  $x$  to generate combinations for. The threads do not need to communicate with each other, avoiding huge overhead especially for large values of  $n$ .

### 3.1.2 Chunk Size

The equations from the section above implicitly defines the chunk size so that each thread gets assigned roughly the same number  $x_i$ 's to work on. It is essentially the number of times task distributions occur in the load balance algorithm, which is roughly  $(n - m + 1)/nth$ .

### 3.1.3 Scheduling Policies

The load balancing algorithm described in Section 3.1.1 implements a static scheduling policy since the threads are assigned specific tasks and only work on the tasks assigned to them. For the purposes of timing comparisons, additional optional arguments `sched` and `chunksize` were added to the function in order to test the speed of the program for the

other scheduling policies, namely dynamic and guided. If both arguments are not provided (or set as NULL), the program defaults to the static scheduling policy using the load balancing algorithm. If either dynamic or guided is set, the program uses OpenMP's built-in `schedule` clause to distribute the tasks to the threads. If the chunksize is not provided, the program sets it to the default chunk size of 1.

We performed various tests using the three scheduling policies and experimented with different chunk sizes (for dynamic and guided). It was clear that the load balancing algorithm using static scheduling was the fastest. It eliminates the communication overhead of the dynamic and guided scheduling policies since the threads don't have to repeatedly access the task farm for work, and idleness is still reduced because of the way the tasks are distributed among the threads.

Hence, the static method is the one compared to the original function in the timing comparisons.

### 3.1.4 Cache and Other Performance Tuning Considerations

The following performance tuning applies only to the load balancing algorithm described in 3.1.1.

- Avoiding instances of false sharing: In order to avoid instances of false sharing (and the resulting cache coherence overheads), we avoided having shared data that could be modified by multiple threads. The data frequently accessed and shared by the threads in the program, such as the input and position vectors are read-only data. Hence, cache lines are not invalidated whenever data from these vectors are accessed. The main shared data structure that is modified by all threads is the output matrix `retmat`, but the threads never have to perform any read on this data.
- Synchronization overhead: The threads are made to work as independent of each other as possible. The only synchronization occurs implicitly at the end of the `parallel` pragma.
- Avoidance of critical sections and barriers

### 3.1.5 Comparative Analysis

The following input sizes were used for comparing the speeds of the source code and our OpenMP implementation using 8 threads. Tests were done using `system.time(combn(sample(1:n), m))`.

```
> nCm(100, 2, 0.100000000000000002)
[1] 4950
> nCm(300, 2, 0.100000000000000002)
[1] 44850
> nCm(100, 3, 0.100000000000000002)
[1] 161700
```


```

> nCm(150, 3, 0.100000000000000002)
[1] 551300
> nCm(200, 3, 0.100000000000000002)
[1] 1313400
> nCm(250, 3, 0.100000000000000002)
[1] 2573000
> nCm(300, 3, 0.100000000000000002)
[1] 4455100

```

The elapsed time recorded for each test is the average of three trials. The OpenMP parallelization massively improved the performance of the function. The improvement becomes increasingly invaluable as the input size increases (e.g. `nCm(100, 5)`, the OpenMP implementation still manages to produce an output under 5 seconds, while the original function takes several minutes).

The following plot illustrates the differences in the speeds of the two implementations.



## 3.2 Snow

### 3.2.1 The Issue of Overhead

After running tests using small inputs, it was observed that parallelization using Snow more often than not took longer than running the CRAN `combn()` (sequentially). This is because of the communication overhead. The communication between the nodes in the cluster took more time than the actual computations in the function. If the jobs sent to the worker nodes are not relatively computationally extensive, the overhead of communicating

ends up deteriorating the performance.

### 3.2.2 Reducing Network Overhead

In order to reduce network overhead and improve the performance, the code was written so that the nodes will do long calculations and less communications. The same load balancing algorithm as in Section 3.1.1 was used to distribute the tasks to each nodes.

### 3.2.3 Comparative Analysis

The same input sizes and function arguments as in Section 3.1.5 were used for testing. For the Snow tests, 8 clusters were used. The following plot illustrates the differences in speeds between the Snow and CRAN implementations.

## 3.3 Thrust

# 4 Timing Comparisons

Test Cases:

```
Test 1: system.time(combn(sample(1:100), 2, NULL, TRUE))
Test 2: system.time(combn(sample(1:300), 2, NULL, TRUE))
Test 3: system.time(combn(sample(1:100), 3, NULL, TRUE))
Test 4: system.time(combn(sample(1:150), 3, NULL, TRUE))
Test 5: system.time(combn(sample(1:200), 3, NULL, TRUE))
Test 6: system.time(combn(sample(1:250), 3, NULL, TRUE))
Test 7: system.time(combn(sample(1:300), 3, NULL, TRUE))
Test 8: system.time(combn(sample(1:120), 3, fun, FALSE))
Test 9: system.time(combn(sample(1:260), 2, NULL, TRUE))
Test 10: system.time(combn(sample(1:180), 2, NULL, TRUE))
Test 11: system.time(combn(sample(1:50), 4, NULL, TRUE))
Test 12: system.time(combn(sample(1:110), 2, fun, TRUE))
Test 13: system.time(combn(sample(1:300), 2, NULL, FALSE))
Test 14: system.time(combn(sample(1:130), 2, fun, FALSE))
Test 15: system.time(combn(sample(1:420), 2, fun, TRUE))
```

```
fun <- function(x) {return(x*x)}
```

The following plot compares all four implementations:

## 5 Conclusions

# A Appendix

```
1 #####
2 # R call function for the OpenMP parallelization of combn() from the CRAN
3 # combinat package: http://cran.r-project.org/web/packages/combinat/index.html
4
5 # Function Arguments:
6 # x <- input vector of integers and/or characters\
7 # m <- number of elements in a combination
8 # fun <- function to apply to the resulting output
9 # simplify <- if TRUE print output as a matrix with m rows and nCm columns
10 # where nCm is the total number of combinations generated
11 # ... <- parameters for fun
12
13 # Helper functions for handling characters in input vector x:
14 # is.letter <- # function to check if there's a char in x
15 # asc <- convert char to ASCII decimal value
16 # chr <- convert decimal value to ASCII character
17 # nCm <- calculating the total number of combinations
18 # (taken directly from R combinat package)
19 # inserted into this file since combinat could not be installed in CSIF
20 #####
21
22 combn <- function(x, m, fun = NULL, simplify = TRUE,
23   sched = NULL, chunksize = NULL, ...)
24 {
25   require(Rcpp)
26   dyn.load("combn-final.so")
27
28   # Input checks taken directly from combn source code
29   if(length(m) > 1) {
30     warning(paste("Argument m has", length(m),
31       "elements: only the first used"))
32     m <- m[1]
33   }
34   if(m < 0)
35     stop("m < 0")
36   if(m == 0)
37     return(if(simplify) vector(mode(x), 0) else list())
38   if(is.numeric(x) && length(x) == 1 && x > 0 && trunc(x) == x)
39     x <- seq(x)
40   n <- length(x)
41   if(n < m)
42     stop("n < m")
43
44   nofun <- is.null(fun)
45   count <- nCm(n, m, 0.100000000000000002)
46
47   # Error checks for the scheduling variables: sched and chunksize
48   # R handles the error when 'sched' is not a string/character vector
49
50   # If sched is provided, then sched must be static, dynamic, guided, or NULL
51   if (!grepl('static', sched) && !grepl('dynamic', sched) && !grepl('guided', sched) &&
52     !is.null(sched)) {
53     stop("Scheduling policy must be static, dynamic, or guided.")
54   }
55   # Set to default values depending on what is/are provided
56   if (is.null(sched) && is.null(chunksize)) {
57     sched <- 'static'
58     chunksize <- 1
59   }
60   else if (!is.null(sched) && is.null(chunksize)) {
61     chunksize <- 1
62   }
63   else if (is.null(sched) && !is.null(chunksize)) { # if sched is provided, but chunk
64     size is not
65     sched <- 'static'
66     warning("'sched' is replaced with default 'static' and 'chunksize' is overridden with
67       default value.")
68   }
```



```

65 }
66
67 # Checks if input vector x has characters
68 # If so, then convert chars to their ASCII decimal values
69 # Operate on the ASCII decimal values for the chars
70 ischarx <- match('TRUE', is.letter(x))
71 if (!is.na(ischarx)) {
72   ischarx_arr <- is.letter(x)
73   for (i in 1:length(charx)) {
74     if (ischarx_arr[i]) {
75       if (length(asc(x[i])) == 1) {
76         x[i] <- asc(x[i])
77       }
78       else {
79         x[i] <- as.character(x[i])
80       }
81     }
82   }
83 }
84 x <- strtoi(x, base=10)
85 }
86
87
88 #Calculate positions for output
89 pos <- vector()
90 temp_n <- n
91 for (i in 1:(n-m+1)) {
92   pos <- c(pos, nCm(temp_n-i, m-1))
93 }
94 temp <- pos[1]
95 pos[1] <- 0
96 for (i in 2:length(pos)) {
97   temp2 <- pos[i]
98   pos[i] <- temp
99   temp <- pos[i] + temp2
100 }
101
102 # Initialize output matrix
103 retmat <- matrix(0, m, count)
104 # Call the function through Rcpp
105 retmat <- .Call("combn", x, m, n, count, sched, chunksize, pos)
106
107 # Convert from ASCII decimal values back to chars if necessary
108 if (!is.na(ischarx)) {
109   for (i in 1:length(retmat)) {
110     if ((as.integer(retmat[i]) >= 97 && as.integer(retmat[i]) <= 122)
111         || (as.integer(retmat[i]) >= 65 && as.integer(retmat[i]) <= 90)) {
112       retmat[i] <- chr(retmat[i]);
113     }
114   }
115 }
116
117 # Apply provided function to the output
118 if (!is.null(fun)) {
119   apply(retmat, 2, fun(...))
120 }
121
122 # Format results
123 if (simplify) {
124   out <- retmat
125 }
126 else {
127   out <- list()
128   for (i in 1:ncol(retmat)) {
129     out <- c(out, list(c(retmat[, i])))
130   }
131 }
132 return(out)
133 }
134

```

```

135 # function to check if there's a char in x
136 is.letter <- function(x) grepl("[[:alpha:]]", x)
137 # convert char to ascii decimal value
138 asc <- function(x) { strtoi(charToRaw(x),16L) }
139 # convert decimal value to ascii character
140 chr <- function(n) { rawToChar(as.raw(n)) }
141
142 # n choose m - calculates the total number of combinations for a given input
143 "nCm"<-
144 function(n, m, tol = 9.9999999999999984e-009)
145 {
146 #   DATE WRITTEN:  7 June 1995                LAST REVISED:  10 July 1995
147 #   AUTHOR:  Scott Chasalow
148 #
149 #   DESCRIPTION:
150 #       Compute the binomial coefficient ("n choose m"), where n is any
151 #       real number and m is any integer. Arguments n and m may be vectors;
152 #       they will be replicated as necessary to have the same length.
153 #
154 #       Argument tol controls rounding of results to integers. If the
155 #       difference between a value and its nearest integer is less than tol,
156 #       the value returned will be rounded to its nearest integer. To turn
157 #       off rounding, use tol = 0. Values of tol greater than the default
158 #       should be used only with great caution, unless you are certain only
159 #       integer values should be returned.
160 #
161 #   REFERENCE:
162 #       Feller (1968) An Introduction to Probability Theory and Its
163 #       Applications, Volume I, 3rd Edition, pp 50, 63.
164 #
165   len <- max(length(n), length(m))
166   out <- numeric(len)
167   n <- rep(n, length = len)
168   m <- rep(m, length = len)
169   mint <- (trunc(m) == m)
170   out[!mint] <- NA
171   out[m == 0] <- 1 # out[mint & (m < 0 | (m > 0 & n == 0))] <- 0
172   whichm <- (mint & m > 0)
173   whichn <- (n < 0)
174   which <- (whichm & whichn)
175   if(any(which)) {
176     nnow <- n[which]
177     mnow <- m[which]
178     out[which] <- ((-1)^mnow) * Recall(mnow - nnow - 1, mnow)
179   }
180   whichn <- (n > 0)
181   nint <- (trunc(n) == n)
182   which <- (whichm & whichn & !nint & n < m)
183   if(any(which)) {
184     nnow <- n[which]
185     mnow <- m[which]
186     foo <- function(j, nn, mm)
187     {
188       n <- nn[j]
189       m <- mm[j]
190       iseq <- seq(n - m + 1, n)
191       negs <- sum(iseq < 0)
192       ((-1)^negs) * exp(sum(log(abs(iseq))) - lgamma(m + 1))
193     }
194     out[which] <- unlist(lapply(seq(along = nnow), foo, nn = nnow,
195       mm = mnow))
196   }
197   which <- (whichm & whichn & n >= m)
198   nnow <- n[which]
199   mnow <- m[which]
200   out[which] <- exp(lgamma(nnow + 1) - lgamma(mnow + 1) - lgamma(nnow -
201     mnow + 1))
202   nna <- !is.na(out)
203   outnow <- out[nna]
204   rout <- round(outnow)

```

```

205     smallldif <- abs(rout - outnow) < tol
206     outnow[smallldif] <- rout[smallldif]
207     out[nna] <- outnow
208     out
209 }

```

```

1  /*****
2  OpenMP (C++) implementation of R's combn() function from the combinat package
3
4  Called from R using .Call() through Rcpp
5  *****/
6
7  #include <Rcpp.h>
8  #include <omp.h>
9
10 using namespace std;
11 using namespace Rcpp;
12
13 // Computes the indices of the next combination to generate
14 // The indices then get mapped to the actual values from the input vector
15 int next_comb(int *comb, int m, int n)
16 {
17     int i = m - 1;
18     ++comb[i];
19
20     while ((i >= 0) && (comb[i] >= n - m + 1 + i)) {
21         --i;
22         ++comb[i];
23     }
24
25     if (comb[0] == 1) {
26         return 0;
27     }
28
29     for (i = i + 1; i < m; ++i) {
30         comb[i] = comb[i - 1] + 1;
31     }
32
33     return 1;
34 }
35
36 RcppExport SEXP combn(SEXP x_, SEXP m_, SEXP n_, SEXP nCm_, SEXP sched_, SEXP chunksize_
    , SEXP pos_, SEXP out)
37 {
38     // Convert SEXP variables to appropriate C++ types
39     NumericVector x(x_); // input vector
40     NumericVector pos(pos_); // position vector for the combinations so that the output is
        sorted
41     int m = as<int>(m_), n = as<int>(n_), nCm = as<int>(nCm_), chunksize = as<int>(
        chunksize_);
42     string sched = as<string>(sched_);
43
44     NumericMatrix retmat(m, nCm);
45
46     // OpenMP schedule clauses
47     if (sched == "dynamic") {
48         omp_set_schedule(omp_sched_dynamic, chunksize);
49     }
50     else if (sched == "guided") {
51         omp_set_schedule(omp_sched_guided, chunksize);
52     }
53
54     if (sched == "static") { // use the load balancing algorithm
55         #pragma omp parallel
56         {
57             // this thread id, total number of threads, combination indexes array
58             int me, nth, *comb;
59
60             nth = omp_get_num_threads();

```

```

61     me = omp_get_thread_num();
62
63     // array that will hold all of the possible combinations
64     // of size m of the indexes
65     comb = new int[m];
66
67     // initialize comb array
68     for (int i = 0; i < m; ++i) {
69         comb[i] = i;
70     }
71
72     int chunkNum = 1; // the number of chunk that has been distributed
73     int mypos; // variable for the output position
74
75     // each thread gets assign a chunk to work on
76     // each thread will have about the same number of chunks
77     // to work on throughout the lifetime of the program
78     for(int current_x = me; current_x < n-m+1; current_x+=1) {
79         int temp;
80         mypos = pos[current_x];
81         for (int i = 0; i < m; ++i) {
82             temp = comb[i] + current_x;
83             retmat(i, mypos) = x[temp];
84         }
85         mypos++;
86         while(next_comb(comb, m, n-current_x)) {
87             int temp;
88             for (int i = 0; i < m; ++i) {
89                 temp = comb[i] + current_x;
90                 retmat(i, mypos) = x[temp];
91             }
92             mypos++;
93         }
94
95         // reset comb array for the next chunk this thread will work on
96         for(int i = 0; i < m; i++) {
97             comb[i] = i;
98         }
99
100        chunkNum++; // increment chunkNum for the next chunk distribution
101        // determine which element this thread will work on
102        if (chunkNum % 2 == 0) {
103            current_x = current_x + 2 * (nth - me - 1);
104        }
105        else {
106            current_x = current_x + 2 * me;
107        }
108    }
109 }
110 }
111 else { // dynamic or guided
112     int mypos;
113     #pragma omp parallel
114     {
115         int *comb = new int[m];
116         for (int i = 0; i < m; ++i) {
117             comb[i] = i;
118         }
119
120         #pragma omp for schedule(runtime)
121         for(int current_x = 0; current_x < (n - m + 1); current_x++) {
122             int temp;
123             mypos = pos[current_x];
124             for (int i = 0; i < m; ++i) {
125                 temp = comb[i] + current_x;
126                 retmat(i, mypos) = x[temp];
127             }
128             mypos++;
129             while(next_comb(comb, m, n-current_x)) {
130                 int temp;

```

```
131         for (int i = 0; i < m; ++i) {
132             temp = comb[i] + current_x;
133             retmat(i, mypos) = x[temp];
134         }
135         mypos++;
136     }
137
138     for(int i = 0; i < m; i++) {
139         comb[i] = i;
140     }
141 }
142 }
143 }
144 return retmat;
145 }
```

## References

- [1] Scott Chasalow *Package ‘combinat’*  
<http://cran.r-project.org/web/packages/combinat/combinat.pdf>
- [2] Junior Barrera, Alfredo Goldman, and Martha Torres *A Parallel Algorithm for Enumerating Combinations* <http://www.ime.usp.br/~gold/ipp03v3.pdf>